

# **Similarity Based Learning**

## **k-Nearest Neighbour**

**Sarah Jane Delany**

**Slides adapted from ML for PDA book**

# Big Idea

- *Looking at what has worked well in the past and make the same (or similar) predictions*

- In 1798, Lieutenant-Colonel David Collins of HMS Calcutta was exploring in NSW when one of his sailors saw a strange animal....

|   | Grrrh! |  |  | Score |
|---|--------|---|---|-------|
|    | ✓      | ✗   | ✗   | 1     |
|   | ✗      | ✓   | ✗   | 1     |
|  | ✗      | ✓   | ✓   | 2     |

# Fundamentals of Similarity Learning


---

- **Feature space** - a  $D$ -dimensional coordinate system used to represent the descriptive features of the instances in the training data, with one axis for each feature
- **Similarity metrics** - the distance between instances in the feature space is a measure of the similarity of the instances

# Feature Space

- Example: 20 college athletes and whether they were drafted by a professional team.
  - Descriptive features are Speed & Agility => 2-d feature space

| ID | Speed | Agility | Draft | ID | Speed | Agility | Draft |
|----|-------|---------|-------|----|-------|---------|-------|
| 1  | 2.50  | 6.00    | No    | 11 | 2.00  | 2.00    | No    |
| 2  | 3.75  | 8.00    | No    | 12 | 5.00  | 2.50    | No    |
| 3  | 2.25  | 5.50    | No    | 13 | 8.25  | 8.50    | No    |
| 4  | 3.25  | 8.25    | No    | 14 | 5.75  | 8.75    | Yes   |
| 5  | 2.75  | 7.50    | No    | 15 | 4.75  | 6.25    | Yes   |
| 6  | 4.50  | 5.00    | No    | 16 | 5.50  | 6.75    | Yes   |
| 7  | 3.50  | 5.25    | No    | 17 | 5.25  | 9.50    | Yes   |
| 8  | 3.00  | 3.25    | No    | 18 | 7.00  | 4.25    | Yes   |
| 9  | 4.00  | 4.00    | No    | 19 | 7.50  | 8.00    | Yes   |
| 10 | 4.25  | 3.75    | No    | 20 | 7.25  | 5.75    | Yes   |




# Measuring Similarity

- Similarity between instances is measured by the distance between the instances
  - Many distance metrics - no 'best' measure => problem dependent

| ID | Speed | Agility | Draft | ID | Speed | Agility | Draft |
|----|-------|---------|-------|----|-------|---------|-------|
| 1  | 2.50  | 6.00    | No    | 11 | 2.00  | 2.00    | No    |
| 2  | 3.75  | 8.00    | No    | 12 | 5.00  | 2.50    | No    |
| 3  | 2.25  | 5.50    | No    | 13 | 8.25  | 8.50    | No    |
| 4  | 3.25  | 8.25    | No    | 14 | 5.75  | 8.75    | Yes   |
| 5  | 2.75  | 7.50    | No    | 15 | 4.75  | 6.25    | Yes   |
| 6  | 4.50  | 5.00    | No    | 16 | 5.50  | 6.75    | Yes   |
| 7  | 3.50  | 5.25    | No    | 17 | 5.25  | 9.50    | Yes   |
| 8  | 3.00  | 3.25    | No    | 18 | 7.00  | 4.25    | Yes   |
| 9  | 4.00  | 4.00    | No    | 19 | 7.50  | 8.00    | Yes   |
| 10 | 4.25  | 3.75    | No    | 20 | 7.25  | 5.75    | Yes   |

Athletes 4 and 5 are close to each other,  
low distance (high similarity)

Athletes 10 and 19 are far from  
each other, high distance (low similarity)



# Distance Metrics

---

- Mathematically a metric must conform to the following four criteria:

1. **Non-negativity**:  $metric(\mathbf{a}, \mathbf{b}) \geq 0$
2. **Identity**:  $metric(\mathbf{a}, \mathbf{b}) = 0 \iff \mathbf{a} = \mathbf{b}$
3. **Symmetry**:  $metric(\mathbf{a}, \mathbf{b}) = metric(\mathbf{b}, \mathbf{a})$
4. **Triangular Inequality**:  
 $metric(\mathbf{a}, \mathbf{b}) \leq metric(\mathbf{a}, \mathbf{c}) + metric(\mathbf{b}, \mathbf{c})$

- Most common distance metric is **Euclidean distance** which computes the length of a straight line between two points,  $\mathbf{a}$  and  $\mathbf{b}$ :

$$Euclidean(\mathbf{a}, \mathbf{b}) = \sqrt{\sum_{i=1}^m (\mathbf{a}[i] - \mathbf{b}[i])^2}$$


# Euclidean Distance

$$Euclidean(\mathbf{a}, \mathbf{b}) = \sqrt{\sum_{i=1}^m (\mathbf{a}[i] - \mathbf{b}[i])^2}$$

$$\begin{aligned} Euclidean(4, 5) &= \sqrt{(3.25 - 2.75)^2 + (8.25 - 7.50)^2} \\ &= \sqrt{(0.25)^2 + (0.75)^2} \\ &= \sqrt{0.625} = 0.7906 \end{aligned}$$

$$Euclidean(10, 19) = ???$$

**Note: Distance and similarity have an inverse relationship**





|    | Speed | Agility |
|----|-------|---------|
| 4  | 3.25  | 8.25    |
| 5  | 2.75  | 7.50    |
| 10 | 4.25  | 3.75    |
| 19 | 7.50  | 8.00    |

# Manhattan Distance

- Another, less well known distance measure is the **Manhattan** distance or *taxi-cab* distance

$$Manhattan(\mathbf{a}, \mathbf{b}) = \sum_{i=1}^m abs(\mathbf{a}[i] - \mathbf{b}[i])$$





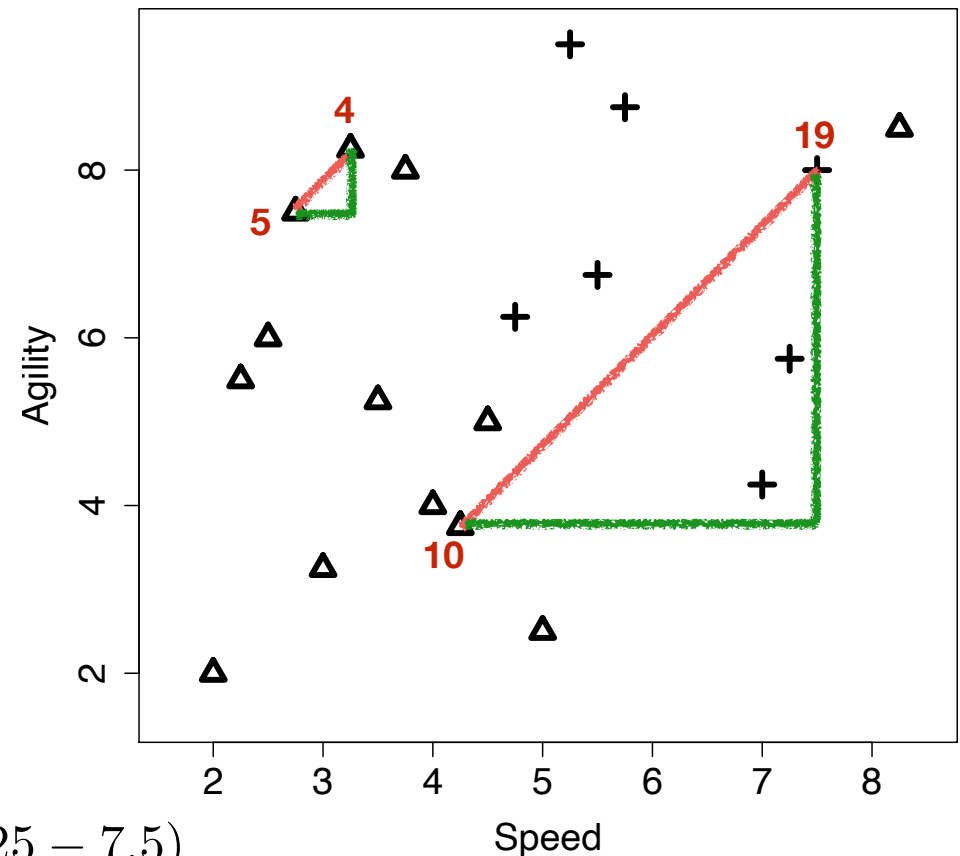
# Manhattan Distance

|    | Speed | Agility |
|----|-------|---------|
| 4  | 3.25  | 8.25    |
| 5  | 2.75  | 7.50    |
| 10 | 4.25  | 3.75    |
| 19 | 7.50  | 8.00    |

$$Manhattan(\mathbf{a}, \mathbf{b}) = \sum_{i=1}^m abs(\mathbf{a}[i] - \mathbf{b}[i])$$

$$\begin{aligned} Manhattan(4, 5) &= abs(3.25 - 2.75) + abs(8.25 - 7.5) \\ &= 0.5 + 0.75 = 1.25 \end{aligned}$$

$$Manhattan(10, 19) = ??$$



Note: Manhattan has a slight computational advantage over Euclidean

# Minkowski Distance

---

- Euclidean and Manhattan distances are special cases of **Minkowski** distance:

$$Minkowski(\mathbf{a}, \mathbf{b}) = \left( \sum_{i=1}^m abs(\mathbf{a}[i] - \mathbf{b}[i])^p \right)^{\frac{1}{p}}$$

- Different values of parameter  $p$  result in different distance measures
  - Manhattan distance for  $p = 1$
  - Euclidean distance for  $p = 2$
- The larger the value of  $p$  the more emphasis is placed on features with large differences in values because these differences are raised to the power of  $p$

# What about non numeric data?

---

- **Binary:** Takes only two values - a boolean True/False decision  
e.g. married={True,False}, test\_result={Pass,Fail}
- **Categorical (Nominal):** A feature that takes values from a finite set of values, with no intrinsic ordering to the values  
e.g. blood\_group={A,B,AB,O}, nationality={French,Irish,Italian}
- **Ordinal:** Similar to a categorical variable, but there is a clear ordering of the variables.  
e.g. grade={A,B,C,D,E,F}, dosage={Low,Medium,High}
- **Interval:** Values that allow ordering and subtraction, but do not allow other arithmetic operations  
e.g date, time

# Categorical Data

- Overlap Difference (feature level):**

Simplest distance measure. Returns 0 if the two values for a feature are equal and 1 otherwise

| <i>Athlete</i> | <i>Gender</i> | <i>Nationality</i> |
|----------------|---------------|--------------------|
| x1             | Female        | Irish              |
| x2             | Male          | Irish              |
| x3             | Male          | Italian            |

For feature  
*Gender*

$$d_g(x1, x2) = 1$$

$$d_g(x1, x3) = 1$$

$$d_g(x2, x3) = 0$$

For feature  
*Nationality*

$$d_n(x1, x2) = 0$$

$$d_n(x1, x3) = 1$$

$$d_n(x2, x3) = 1$$

- Hamming distance:** Distance metric for instance represented with categorical data only, = the sum of the overlap differences across all features - i.e. number of features on which two examples disagree.

$$d(x1, x2) = 1 + 0 = 1$$

$$d(x1, x3) = 1 + 1 = 2$$

$$d(x2, x3) = 0 + 1 = 1$$

Overlap distance for *Gender* +  
Overlap distance for *Nationality*

# Ordinal Data

---

- For *ordinal features*, calculate the absolute value of the difference between the two positions in the ordered list of possible values.

e.g. Ordinal Feature *Dosage*:  
 $\{\text{Low}, \text{Medium}, \text{High}\} = \{1, 2, 3\}$

$\text{diff}(\text{Low}, \text{High}) = |1-3| = 2$   
 $\text{diff}(\text{Medium}, \text{Low}) = |2-1| = 1$   
 $\text{diff}(\text{High}, \text{High}) = |3-3| = 0$

# Heterogeneous Distance Measures

---

- In many datasets, the features associated with instances will have different types (e.g. continuous, categorical, ordinal etc).

- **Local distance function:** Measure the distance between two instances based on a single feature.

| <i>Athlete</i> | <i>Speed</i> | <i>Agility</i> | <i>Gender</i> | <i>Nationality</i> |
|----------------|--------------|----------------|---------------|--------------------|
| <b>x1</b>      | 2.50         | 6.00           | Female        | Irish              |
| <b>x2</b>      | 3.75         | 8.00           | Male          | Irish              |
| <b>x3</b>      | 2.25         | 5.50           | Male          | Italian            |

- e.g. distance between **x1** and **x2** in terms of *Speed*?
  - e.g. distance between **x1** and **x3** in terms of *Gender*?
  - e.g. distance between **x1** and **x2** in terms of *Nationality*?
- **Global distance function:** Measure the distance between two instances based on the combination of the local distances across all features.

# Heterogeneous Distance Functions

- We can create a **global measure** from different local distance functions, using an appropriate function for each feature.

| <i>Athlete</i> | <i>Speed</i> | <i>Agility</i> | <i>Gender</i> | <i>Nationality</i> |
|----------------|--------------|----------------|---------------|--------------------|
| <b>x1</b>      | 2.50         | 6.00           | Female        | Irish              |
| <b>x2</b>      | 3.75         | 8.00           | Male          | Irish              |
| <b>x3</b>      | 2.25         | 5.50           | Male          | Italian            |

Use absolute difference for continuous features *Speed & Agility*

Use overlap for categorical features *Gender & Nationality*

$$d(x1, x2) = 1.25 + 2.0 + 1 + 0 = 4.25$$

$$d(x1, x3) = 0.25 + 0.5 + 1 + 1 = 2.75$$

$$d(x2, x3) = 1.5 + 2.5 + 0 + 1 = 5.0$$

Global distance calculated as sum over individual local distances

- Often domain expertise is required to choose an appropriate distance measure for a particular dataset.

# Nearest Neighbour Algorithm (k-NN)

---

- Have a set of training instances and a query to be classified
- Steps:
  - Iterate across the training instances in memory and find the instance(s) that is/are the most similar (shortest distance) to the query instance in the feature space
  - Make a prediction for the query instance based on the target values of the nearest neighbour(s) of the query instance
- **1-NN** - use the most similar/closest training instance
- **k-NN** - use the k most similar/closest training instances



# Example


---

| ID | Speed | Agility | Draft | ID | Speed | Agility | Draft |
|----|-------|---------|-------|----|-------|---------|-------|
| 1  | 2.50  | 6.00    | No    | 11 | 2.00  | 2.00    | No    |
| 2  | 3.75  | 8.00    | No    | 12 | 5.00  | 2.50    | No    |
| 3  | 2.25  | 5.50    | No    | 13 | 8.25  | 8.50    | No    |
| 4  | 3.25  | 8.25    | No    | 14 | 5.75  | 8.75    | Yes   |
| 5  | 2.75  | 7.50    | No    | 15 | 4.75  | 6.25    | Yes   |
| 6  | 4.50  | 5.00    | No    | 16 | 5.50  | 6.75    | Yes   |
| 7  | 3.50  | 5.25    | No    | 17 | 5.25  | 9.50    | Yes   |
| 8  | 3.00  | 3.25    | No    | 18 | 7.00  | 4.25    | Yes   |
| 9  | 4.00  | 4.00    | No    | 19 | 7.50  | 8.00    | Yes   |
| 10 | 4.25  | 3.75    | No    | 20 | 7.25  | 5.75    | Yes   |

- Should we draft an athlete who with the following profile?

Speed = 6.75; Agility = 3

# Example




Triangle: No  
Plus: Yes

# Example: 1-NN

- Compute distance (using selected distance measure) of query to each training instance
  - Rank all instances based on calculated distance
  - See which is the *closest* to query
- Remember: lowest distance = highest similarity

| ID | SPEED | AGILITY | DRAFT | Dist. | ID | SPEED | AGILITY | DRAFT | Dist. |
|----|-------|---------|-------|-------|----|-------|---------|-------|-------|
| 18 | 7.00  | 4.25    | yes   | 1.27  | 11 | 2.00  | 2.00    | no    | 4.85  |
| 12 | 5.00  | 2.50    | no    | 1.82  | 19 | 7.50  | 8.00    | yes   | 5.06  |
| 10 | 4.25  | 3.75    | no    | 2.61  | 3  | 2.25  | 5.50    | no    | 5.15  |
| 20 | 7.25  | 5.75    | yes   | 2.80  | 1  | 2.50  | 6.00    | no    | 5.20  |
| 9  | 4.00  | 4.00    | no    | 2.93  | 13 | 8.25  | 8.50    | no    | 5.70  |
| 6  | 4.50  | 5.00    | no    | 3.01  | 2  | 3.75  | 8.00    | no    | 5.83  |
| 8  | 3.00  | 3.25    | no    | 3.76  | 14 | 5.75  | 8.75    | yes   | 5.84  |
| 15 | 4.75  | 6.25    | yes   | 3.82  | 5  | 2.75  | 7.50    | no    | 6.02  |
| 7  | 3.50  | 5.25    | no    | 3.95  | 4  | 3.25  | 8.25    | no    | 6.31  |
| 16 | 5.50  | 6.75    | yes   | 3.95  | 17 | 5.25  | 9.50    | yes   | 6.67  |

# Example: 3-NN



Triangle: No  
Plus: Yes

# Example: 3-NN

- Compute distance of query to each training instance
- Rank all instances based on distance
- Choose the  $k$  nearest neighbours
- Determine predicted target by **majority vote** of target of the  $k$  nearest neighbours

| ID | SPEED | AGILITY | DRAFT | Dist. | ID | SPEED | AGILITY | DRAFT | Dist. |
|----|-------|---------|-------|-------|----|-------|---------|-------|-------|
| 18 | 7.00  | 4.25    | yes   | 1.27  | 11 | 2.00  | 2.00    | no    | 4.85  |
| 12 | 5.00  | 2.50    | no    | 1.82  | 19 | 7.50  | 8.00    | yes   | 5.06  |
| 10 | 4.25  | 3.75    | no    | 2.61  | 3  | 2.25  | 5.50    | no    | 5.15  |
| 20 | 7.25  | 5.75    | yes   | 2.80  | 1  | 2.50  | 6.00    | no    | 5.20  |
| 9  | 4.00  | 4.00    | no    | 2.93  | 13 | 8.25  | 8.50    | no    | 5.70  |
| 6  | 4.50  | 5.00    | no    | 3.01  | 2  | 3.75  | 8.00    | no    | 5.83  |
| 8  | 3.00  | 3.25    | no    | 3.76  | 14 | 5.75  | 8.75    | yes   | 5.84  |
| 15 | 4.75  | 6.25    | yes   | 3.82  | 5  | 2.75  | 7.50    | no    | 6.02  |
| 7  | 3.50  | 5.25    | no    | 3.95  | 4  | 3.25  | 8.25    | no    | 6.31  |
| 16 | 5.50  | 6.75    | yes   | 3.95  | 17 | 5.25  | 9.50    | yes   | 6.67  |

What about 4-NN?

# Example: 4-NN

| ID | SPEED | AGILITY | DRAFT | Dist. | ID | SPEED | AGILITY | DRAFT | Dist. |
|----|-------|---------|-------|-------|----|-------|---------|-------|-------|
| 18 | 7.00  | 4.25    | yes   | 1.27  | 11 | 2.00  | 2.00    | no    | 4.85  |
| 12 | 5.00  | 2.50    | no    | 1.82  | 19 | 7.50  | 8.00    | yes   | 5.06  |
| 10 | 4.25  | 3.75    | no    | 2.61  | 3  | 2.25  | 5.50    | no    | 5.15  |
| 20 | 7.25  | 5.75    | yes   | 2.80  | 1  | 2.50  | 6.00    | no    | 5.20  |
| 9  | 4.00  | 4.00    | no    | 2.93  | 13 | 8.25  | 8.50    | no    | 5.70  |
| 6  | 4.50  | 5.00    | no    | 3.01  | 2  | 3.75  | 8.00    | no    | 5.83  |
| 8  | 3.00  | 3.25    | no    | 3.76  | 14 | 5.75  | 8.75    | yes   | 5.84  |
| 15 | 4.75  | 6.25    | yes   | 3.82  | 5  | 2.75  | 7.50    | no    | 6.02  |
| 7  | 3.50  | 5.25    | no    | 3.95  | 4  | 3.25  | 8.25    | no    | 6.31  |
| 16 | 5.50  | 6.75    | yes   | 3.95  | 17 | 5.25  | 9.50    | yes   | 6.67  |

- Can break ties
  - Randomly
  - Based on the sum of the nearest neighbour distances for each target class

# *k*-NN algorithm

---

- *k* nearest neighbours algorithm predicts the target class with the majority vote from the set of *k* nearest neighbours to the query *q*:

$$\mathbb{M}_k(\mathbf{q}) = \operatorname{argmax}_{c \in \text{classes}(t)} \sum_{i=1}^k \delta_{t_i, c}$$

*classes*(*t*) is the set of target classes

*t<sub>i</sub>* is the target class for instance *i*

$\delta_{i,j}$  is Kronecker's delta  $\delta_{i,j} = \begin{cases} 1 & \text{when } i = j \\ 0 & \text{when } i \neq j. \end{cases}$

# Weighted $k$ -NN algorithm

---

- In distance weighted  $k$  Nearest Neighbour algorithm some neighbours get higher weight than others

$$\mathbb{M}_k(\mathbf{q}) = \arg \max_{c \in \text{classes}(t)} \sum_{i=1}^k w_i \times \delta_{t_i, c}$$


- Instead of a binary vote of 1 for the class of the neighbour, closest neighbours (those that are most similar to the query) get a higher weighting when deciding the prediction for the query
- Remember: similarity = inverse of distance

$$w_i = \frac{1}{\text{dist}(\mathbf{q}, \mathbf{d}_i)}$$



# Tuning for $k$

- A simple 1-NN classifier is easy to implement. But it will be susceptible to “noise” in the data. A misclassification will occur every time a single noisy example is retrieved.
- We might decide to vary the neighbourhood size parameter  $k$  to improve the predictive performance of  $k$ -NN.
- Choosing between different settings of an algorithm is often referred to as *hyperparameter tuning* or *model selection*.
- Using a larger  $k$  (e.g.  $k > 2$ ) can sometimes make the classifier more robust and overcome this problem.
- But when  $k$  is large ( $k \rightarrow N$ ) and classes are *unbalanced*, we always predict the majority class.




# Data Normalisation

- Consider the dataset listing salary and age information for customers and whether or not they purchased a pension plan.

| ID | Salary | Age | Purchased |
|----|--------|-----|-----------|
| 1  | 53700  | 41  | No        |
| 2  | 65300  | 37  | No        |
| 3  | 48900  | 45  | Yes       |
| 4  | 64800  | 49  | Yes       |
| 5  | 44200  | 30  | No        |
| 6  | 55900  | 57  | Yes       |
| 7  | 48600  | 26  | No        |
| 8  | 72800  | 60  | Yes       |
| 9  | 45300  | 34  | No        |
| 10 | 73200  | 52  | Yes       |

What should the marketing dept expect for customer aged 35, with salary of 56,000?



# Data Normalisation

- Calculate distance using Euclidean distance:

| ID | Salary | Age | Purch. | Salary and Age |       | Salary Only |       | Age Only |       |
|----|--------|-----|--------|----------------|-------|-------------|-------|----------|-------|
|    |        |     |        | Dist.          | Rank. | Dist.       | Rank. | Dist.    | Rank. |
| 1  | 53700  | 41  | No     | 2300.0078      | 2     | 2300        | 2     | 6        | 4     |
| 2  | 65300  | 37  | No     | 9300.0002      | 6     | 9300        | 6     | 2        | 2     |
| 3  | 48900  | 45  | Yes    | 7100.0070      | 3     | 7100        | 3     | 10       | 6     |
| 4  | 64800  | 49  | Yes    | 8800.0111      | 5     | 8800        | 5     | 14       | 7     |
| 5  | 44200  | 30  | No     | 11800.0011     | 8     | 11800       | 8     | 5        | 3     |
| 6  | 55900  | 57  | Yes    | 102.3914       | 1     | 100         | 1     | 22       | 9     |
| 7  | 48600  | 26  | No     | 7400.0055      | 4     | 7400        | 4     | 9        | 5     |
| 8  | 72800  | 60  | Yes    | 16800.0186     | 9     | 16800       | 9     | 25       | 10    |
| 9  | 45300  | 34  | No     | 10700.0000     | 7     | 10700       | 7     | 1        | 1     |
| 10 | 73200  | 52  | Yes    | 17200.0084     | 10    | 17200       | 10    | 17       | 8     |


- Salary feature dominates the computation of distance, age feature is virtually ignored
- Due to the larger range in the Salary feature

# Data Normalisation

- Use **min-max** range normalisation to rescale values to the range of [0,1]

$$a'_i = \frac{a_i - \min(a)}{\max(a) - \min(a)} \times (high - low) + low$$

$$= \frac{a_i - \min(a)}{\max(a) - \min(a)}$$



| ID | Normalized Dataset |        |        | Salary and Age |      | Salary Only |      | Age Only |      |
|----|--------------------|--------|--------|----------------|------|-------------|------|----------|------|
|    | Salary             | Age    | Purch. | Dist.          | Rank | Dist.       | Rank | Dist.    | Rank |
| 1  | 0.3276             | 0.4412 | No     | 0.1935         | 1    | 0.0793      | 2    | 0.17647  | 4    |
| 2  | 0.7276             | 0.3235 | No     | 0.3260         | 2    | 0.3207      | 6    | 0.05882  | 2    |
| 3  | 0.1621             | 0.5588 | Yes    | 0.3827         | 5    | 0.2448      | 3    | 0.29412  | 6    |
| 4  | 0.7103             | 0.6765 | Yes    | 0.5115         | 7    | 0.3034      | 5    | 0.41176  | 7    |
| 5  | 0.0000             | 0.1176 | No     | 0.4327         | 6    | 0.4069      | 8    | 0.14706  | 3    |
| 6  | 0.4034             | 0.9118 | Yes    | 0.6471         | 8    | 0.0034      | 1    | 0.64706  | 9    |
| 7  | 0.1517             | 0.0000 | No     | 0.3677         | 3    | 0.2552      | 4    | 0.26471  | 5    |
| 8  | 0.9862             | 1.0000 | Yes    | 0.9361         | 10   | 0.5793      | 9    | 0.73529  | 10   |
| 9  | 0.0379             | 0.2353 | No     | 0.3701         | 4    | 0.3690      | 7    | 0.02941  | 1    |
| 10 | 1.0000             | 0.7647 | Yes    | 0.7757         | 9    | 0.5931      | 10   | 0.50000  | 8    |

# Predicting Continuous Targets

---

- $k$  nearest neighbours algorithm predicts the average target value of  $k$  nearest neighbours to the query  $q$ :

$$\mathbb{M}_k(\mathbf{q}) = \frac{1}{k} \sum_{i=1}^k t_i$$

$t_i$  is the target feature value for instance  $i$

# *k*-NN Algorithm

---

- Similarity based learning attempts to mimic a very human way of reasoning - this makes them more **explainable** - easy to interpret and understand
  - Can give people more confidence in the model
- **Lazy learning technique**: Model is built at classification time rather than training time
  - Uses a set of local training instances to classify each query
  - Appropriate for heterogeneous data
  - Computationally more expensive as the number of training instances becomes larger
- Easy to add new instances to training data for re-training to handle **concept drift**
- Important to normalise data (for all prediction algorithms)